




Overview

TensorFlow began its life in 2011 as DisBelief, an internal, closed source project at Google. DisBelief was a machine learning system that employed deep learning neural networks. This system morphed into TensorFlow, which was released to the developer community under an Apache 2.0 open source license, on November 9, 2015. Version 1.0.0 made its appearance on February 11, 2017. There have been a number of point releases since then that have incorporated a wealth of new features.

TensorFlow quickly rose to become the most widely used machine learning library today. And not without reason.



Google released the TensorFlow 2.0 alpha version in March 2019. The official 2.0 release is expected in Q2 2019. TensorFlow 2.0 promises simplicity

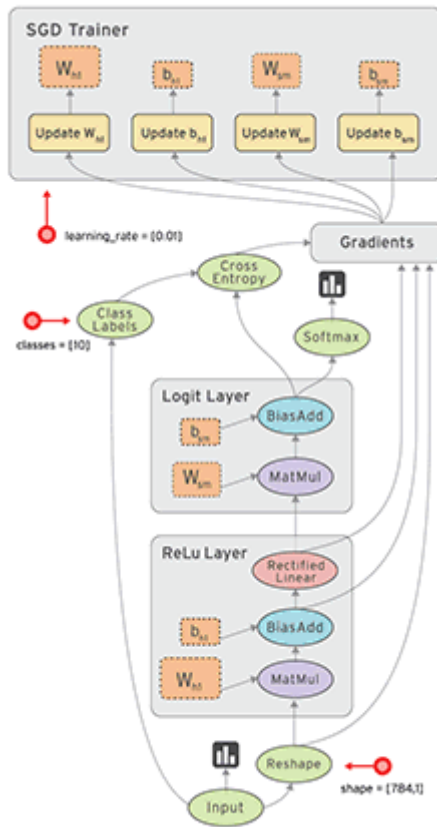
TensorFlow Ecosystem

- Language Support
- Computation Graph
- Gradients
- CUDA
- Multi-GPU Support
- Eager Execution (> 1.5, > 1.7)
- Debugger
- TensorBoard
- TPU Support
- TensorRT
- `tf.contrib`
- TensorFlow Hub

TensorFlow 1.x [`tf.contrib`]

<code>__pycache__/</code>	<code>deprecated/</code>	<code>kafka/</code>	<code>nn/</code>	<code>signal/</code>
<code>all_reduce/</code>	<code>distribute/</code>	<code>keras/</code>	<code>opt/</code>	<code>slim/</code>
<code>autograph/</code>	<code>distributions/</code>	<code>kernel_methods/</code>	<code>optimizer_v2/</code>	<code>solvers/</code>
<code>batching/</code>	<code>eager/</code>	<code>kinesis/</code>	<code>periodic_resample/</code>	<code>sparsemax/</code>
<code>bayesflow/</code>	<code>estimator/</code>	<code>labeled_tensor/</code>	<code>predictor/</code>	<code>specs/</code>
<code>bigtable/</code>	<code>factorization/</code>	<code>layers/</code>	<code>proto/</code>	<code>staging/</code>
<code>boosted_trees/</code>	<code>feature_column/</code>	<code>learn/</code>	<code>quantization/</code>	<code>stat_summarizer/</code>
<code>checkpoint/</code>	<code>ffmpeg/</code>	<code>legacy_seq2seq/</code>	<code>quantize/</code>	<code>stateless/</code>
<code>cloud/</code>	<code>framework/</code>	<code>libsvm/</code>	<code>rate/</code>	<code>summary/</code>
<code>cluster_resolver/</code>	<code>fused_conv/</code>	<code>linear_optimizer/</code>	<code>receptive_field/</code>	<code>tensor_forest/</code>
<code>cmake/</code>	<code>gan/</code>	<code>lite/</code>	<code>recurrent/</code>	<code>tensorboard/</code>
<code>coder/</code>	<code>graph_editor/</code>	<code>lookup/</code>	<code>reduce_slice_ops/</code>	<code>tensorrt/</code>
<code>compiler/</code>	<code>grid_rnn/</code>	<code>losses/</code>	<code>remote_fused_graph/</code>	<code>testing/</code>
<code>constrained_optimization/</code>	<code>hadoop/</code>	<code>memory_stats/</code>	<code>resampler/</code>	<code>text/</code>
<code>copy_graph/</code>	<code>hooks/</code>	<code>meta_graph_transform/</code>	<code>rnn/</code>	<code>tfprof/</code>
<code>crf/</code>	<code>ignite/</code>	<code>metrics/</code>	<code>rpc/</code>	<code>timeseries/</code>
<code>cuda_rnn/</code>	<code>image/</code>	<code>mixed_precision/</code>	<code>saved_model/</code>	<code>tpu/</code>
<code>data/</code>	<code>input_pipeline/</code>	<code>model_pruning/</code>	<code>seq2seq/</code>	<code>training/</code>
<code>decision_trees/</code>	<code>integrate/</code>	<code>nearest_neighbor/</code>	<code>session_bundle/</code>	<code>util/</code>

TensorFlow 1.x [Computation Graph]



TensorFlow 1.x Vs. 2.0 [Eager Mode]

```
[2] 1 import tensorflow as tf
    2
    3 # define the inputs
    4 x = tf.placeholder(tf.float32)
    5 y = tf.placeholder(tf.float32)
    6
    7 # define the graph
    8 g_mean = tf.sqrt(x * y)
    9
   10 # run the graph
   11 with tf.Session() as session:
   12     result = session.run(g_mean, feed_dict={x: 2, y: 8})
   13     print(result)
```

4.0

```
[2] 1 import tensorflow as tf
    2
    3 # define the inputs
    4 x = 2.0
    5 y = 8.0
    6
    7 # define the graph
    8 g_mean = tf.sqrt(x * y)
    9
   10 # run the graph
   11 tf.print(g_mean)
```

4

TensorFlow 1.x Vs. 2.0 [Graph]

```
[2] 1 import tensorflow as tf
    2
    3 # define the inputs
    4 x = tf.placeholder(tf.float32)
    5 y = tf.placeholder(tf.float32)
    6
    7 # define the graph
    8 g_mean = tf.sqrt(x * y)
    9
   10 # run the graph
   11 with tf.Session() as session:
   12     result = session.run(g_mean, feed_dict={x: 2, y: 8})
   13     print(result)
```

4.0

```
[6] 1 import tensorflow as tf
    2
    3 # define the inputs
    4 x = 2.0
    5 y = 8.0
    6
    7 # define the graph
    8 @tf.function
    9 def geometric_mean(x, y):
   10     g_mean = tf.sqrt(x * y)
   11     return g_mean
   12
   13 # run the graph
   14 g_mean = geometric_mean(x, y)
   15 tf.print(g_mean)
```

4

TensorFlow 1.x Vs. 2.0 [Architecture Layer]

```

1 import tensorflow as tf
2
3
4 def preprocess_data(im, label):
5     im = tf.cast(im, tf.float32)
6     im = im / 127.5
7     im = im - 1
8     im = tf.reshape(im, [-1])
9     return im, label
10
11 def data_layer(data_tensor, num_threads=8, prefetch_buffer=100, batch_size=32):
12     with tf.variable_scope("data"):
13         dataset = tf.data.Dataset.from_tensor_slices(data_tensor)
14         dataset = dataset.shuffle(buffer_size=60000).repeat()
15         dataset = dataset.map(preprocess_data, num_parallel_calls=num_threads)
16         dataset = dataset.batch(batch_size)
17         dataset = dataset.prefetch(prefetch_buffer)
18         iterator = dataset.make_one_shot_iterator()
19     return iterator
20
21 def model(input_layer, num_classes=10):
22     with tf.variable_scope("model"):
23         net = tf.layers.dense(input_layer, 512)
24         net = tf.nn.relu(net)
25         net = tf.layers.dense(net, num_classes)
26     return net
27
28 def loss_functions(logits, labels, num_classes=10):
29     with tf.variable_scope("loss"):
30         target_prob = tf.one_hot(labels, num_classes)
31         total_loss = tf.losses.softmax_cross_entropy(target_prob, logits)
32     return total_loss
33
34 def optimizer_func(total_loss, global_step, learning_rate=0.1):
35     with tf.variable_scope("optimizer"):
36         optimizer = tf.train.GradientDescentOptimizer(learning_rate=learning_rate)
37         optimizer = optimizer.minimize(total_loss, global_step=global_step)
38     return optimizer
39
40 def performance_metric(logits, labels):
41     with tf.variable_scope("performance_metric"):
42         preds = tf.argmax(logits, axis=-1)
43         labels = tf.cast(labels, tf.int64)
44         corrects = tf.equal(preds, labels)
45         accuracy = tf.reduce_mean(tf.cast(corrects, tf.float32))
46     return accuracy
47
48 def train(data_tensor):
49     global_step = tf.Variable(1, dtype=tf.int32, trainable=False, name="iter_number")
50
51     # training graph
52     images, labels = data_layer(data_tensor).get_next()
53     logits = model(images)
54     loss = loss_functions(logits, labels)
55     optimizer = optimizer_func(loss, global_step)
56     accuracy = performance_metric(logits, labels)
57
58     # start training
59     num_iter = 10000
60     log_iter = 1000
61     with tf.Session() as sess:
62         sess.run(tf.global_variables_initializer())
63         streaming_loss = 0
64         streaming_accuracy = 0
65
66         for i in range(1, num_iter + 1):
67             _, loss_batch, acc_batch = sess.run([optimizer, loss, accuracy])
68             streaming_loss += loss_batch
69             streaming_accuracy += acc_batch
70             if i % log_iter == 0:
71                 print("Iteration: {}, Streaming loss: {:.2f}, Streaming accuracy: {:.2f}"
72                       .format(i, streaming_loss/log_iter, streaming_accuracy/log_iter))
73             streaming_loss = 0
74             streaming_accuracy = 0
75
76 data_train, data_val = tf.keras.datasets.mnist.load_data()
77 print(data_train[0].shape, data_train[1].shape, data_val[0].shape, data_val[1].shape)
78 train(data_tensor=data_train)

```

```

[ ] 1 import tensorflow as tf
2
3
4 (train_images, train_labels), (test_images, test_labels) = tf.keras.datasets.mnist.load_data()
5
6 model = tf.keras.Sequential([
7     tf.keras.layers.Flatten(input_shape=(28, 28)),
8     tf.keras.layers.Dense(512, activation=tf.nn.relu),
9     tf.keras.layers.Dense(10, activation=tf.nn.softmax),
10 ])
11
12
13 model.compile(optimizer='adam',
14               loss='sparse_categorical_crossentropy',
15               metrics=['accuracy'])
16
17
18 model.fit(train_images, train_labels, epochs=5)

```

TensorFlow 2.0 [Defining Model v1]

```
[ ] 1 import tensorflow as tf
    2
    3
    4 (train_images, train_labels), (test_images, test_labels) = tf.keras.datasets.mnist.load_data()
    5
    6 model = tf.keras.Sequential([
    7     tf.keras.layers.Flatten(input_shape=(28, 28)),
    8     tf.keras.layers.Dense(512, activation=tf.nn.relu),
    9     tf.keras.layers.Dense(10, activation=tf.nn.softmax),
10 ])
11
12
13 model.compile(optimizer='adam',
14               loss='sparse_categorical_crossentropy',
15               metrics=['accuracy'])
16
17
18 model.fit(train_images, train_labels, epochs=5)
```


TensorFlow 2.0 [Defining Model v2]

```
[ ] 1 import tensorflow as tf
    2
    3
    4 (train_images, train_labels), (test_images, test_labels) = tf.keras.datasets.mnist.load_data()
    5
    6 inputs = tf.keras.Input(shape=[28, 28])
    7 net = tf.keras.layers.Flatten()(inputs)
    8 net = tf.keras.layers.Dense(512, activation=tf.nn.relu)(net)
    9 net = tf.keras.layers.Dense(10, activation=tf.nn.softmax)(net)
   10 model = tf.keras.Model(inputs=inputs, outputs=net)
   11
   12 model.compile(optimizer='adam',
   13               loss='sparse_categorical_crossentropy',
   14               metrics=['accuracy'])
   15
   16
   17 model.fit(train_images, train_labels, epochs=5)
```

TensorFlow 2.0 [Defining Model v3]

```
[ ] 1 import tensorflow as tf
    2
    3
    4 train_data, _ = tf.keras.datasets.mnist.load_data()
    5
    6 dataset = tf.data.Dataset.from_tensor_slices(train_data)
    7 dataset = dataset.shuffle(buffer_size=60000)
    8 dataset = dataset.batch(32)
    9
   10 model = tf.keras.Sequential([
   11     tf.keras.layers.Flatten(input_shape=(28, 28)),
   12     tf.keras.layers.Dense(512, activation=tf.nn.relu),
   13     tf.keras.layers.Dense(10, activation=tf.nn.softmax),
   14 ])
   15
   16 model.compile(optimizer='adam',
   17               loss='sparse_categorical_crossentropy',
   18               metrics=['accuracy'])
   19
   20
   21 model.fit(dataset, epochs=5)
```