



Welcome to Session 02 [Convolutional Neural Network


(CNN)]

This is a [Google Colaboratory \(https://colab.research.google.com/notebooks/welcome.ipynb\)](https://colab.research.google.com/notebooks/welcome.ipynb) notebook file. Python programs are run directly in the browser—a great way to learn and use TensorFlow. To follow this tutorial, run the notebook in Google Colab by clicking the button at the top of this page.

1. In Colab, connect to a Python runtime: At the top-right of the menu bar, select *CONNECT*.
2. Run all the notebook code cells: Select *Runtime > Run all*.


Overview

Architecture of a traditional CNN — Convolutional neural networks, also known as CNNs, are a specific type of neural networks that are generally composed of the following layers:



Types of Layer

Convolution layer (CONV) — The convolution layer (CONV) uses filters that perform convolution operations as it is scanning the input I with respect to its dimensions. Its hyperparameters include the filter size F and stride S . The resulting output O is called feature map or activation map.



0	0	0	0	0	0	...
0	156	155	156	158	158	...
0	153	154	157	159	159	...
0	149	151	155	158	159	...
0	146	146	149	153	158	...
0	145	143	143	148	158	...
...

Input Channel #1 (Red)

0	0	0	0	0	0	...
0	167	166	167	169	169	...
0	164	165	168	170	170	...
0	160	162	166	169	170	...
0	156	156	159	163	168	...
0	155	153	153	158	168	...
...

Input Channel #2 (Green)

0	0	0	0	0	0	...
0	163	162	163	165	165	...
0	160	161	164	166	166	...
0	156	158	162	165	166	...
0	155	155	158	162	167	...
0	154	152	152	157	167	...
...

Input Channel #3 (Blue)

-1	-1	1
0	1	-1
0	1	1

Kernel Channel #1



308

+

1	0	0
1	-1	-1
1	0	-1

Kernel Channel #2



-498

+

0	1	1
0	1	0
1	-1	1

Kernel Channel #3



164

+ 1 = -25



Bias = 1

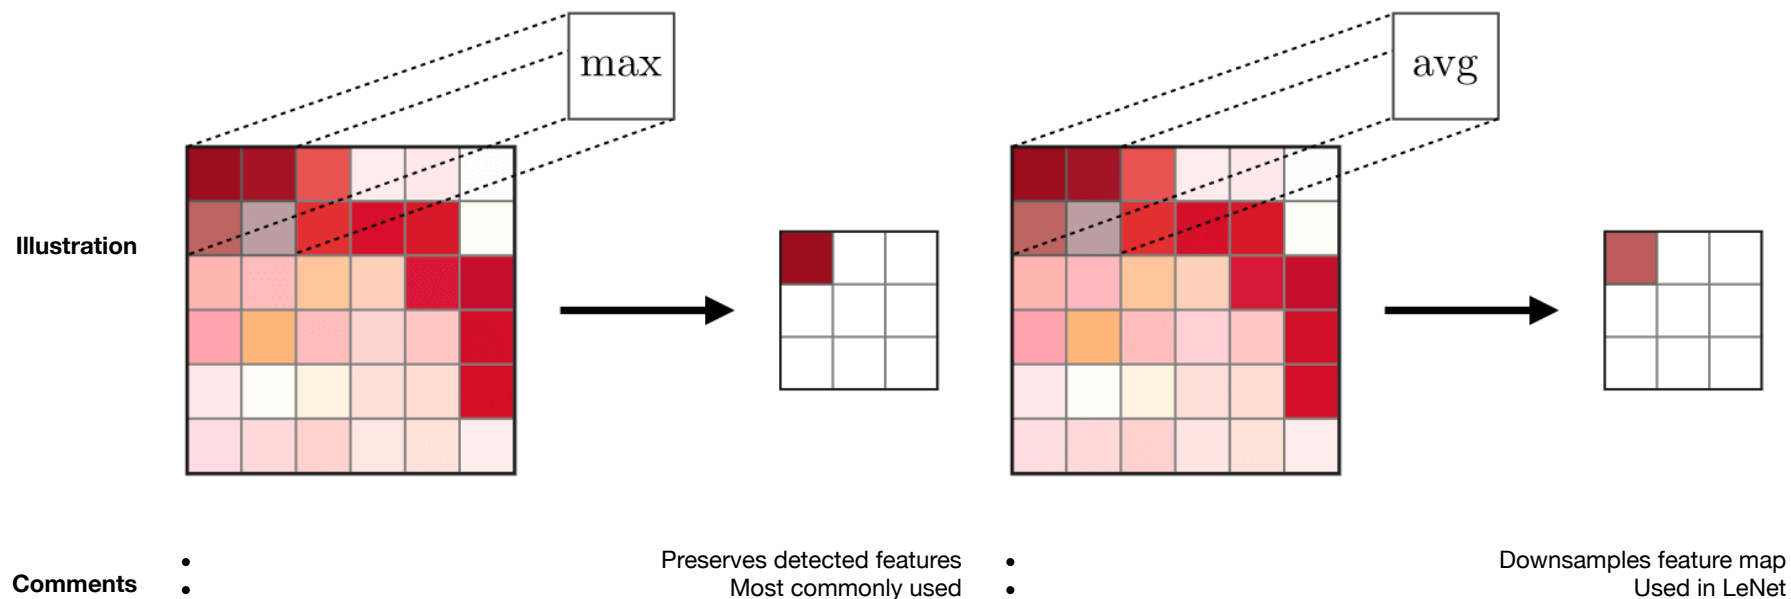
Output

-25				...
				...
				...
				...
...


Remark: the convolution step can be generalized to the 1D and 3D cases as well.

Pooling (POOL) — The pooling layer (POOL) is a downsampling operation, typically applied after a convolution layer, which does some spatial invariance. In particular, max and average pooling are special kinds of pooling where the maximum and average value is taken, respectively.

Type	Max pooling	Average pooling
Purpose	Each pooling operation selects the maximum value of the current view	Each pooling operation averages the values of the current view




Fully Connected (FC) — The fully connected layer (FC) operates on a flattened input where each input is connected to all neurons. If present, FC layers are usually found towards the end of CNN architectures and can be used to optimize objectives such as class scores.



Filter hyperparameters

The convolution layer contains filters for which it is important to know the meaning behind its hyperparameters.

Dimensions of a filter — A filter of size $F \times F$ applied to an input containing C channels is a $F \times F \times C$ volume that performs convolutions on an input of size $I \times I \times C$ and produces an output feature map (also called activation map) of size $O \times O \times 1$.



Remark: the application of K filters of size $F \times F$ results in an output feature map of size $O \times O \times K$.

Stride — For a convolutional or a pooling operation, the stride S denotes the number of pixels by which the window moves after each operation.



Zero-padding — Zero-padding denotes the process of adding P zeroes to each side of the boundaries of the input. This value can either be manually specified or automatically set through one of the three modes detailed below:

Mode	Valid	Same	Full
Illustration			
Purpose	<ul style="list-style-type: none"> No padding Drops last convolution if dimensions do not match 	<ul style="list-style-type: none"> Output size is mathematically convenient Also called 'half' padding 	<ul style="list-style-type: none"> Maximum padding such that end convolutions are applied on the limits of the input Filter 'sees' the input end-to-end

Commonly used activation functions

Rectified Linear Unit — The rectified linear unit layer (ReLU) is an activation function g that is used on all elements of the volume. It aims at introducing non-linearity to the network. Its variants are summarized in the table below:

Implementing The Model

Initialization

Mounting

```
In [0]: # Mounting Gdrive

USE_G_COLAB = True

if USE_G_COLAB:
    from google.colab import drive

    drive.mount('/content/drive', force_remount=True)
```

Go to this URL in a browser: https://accounts.google.com/o/oauth2/auth?client_id=947318989803-6bn6qk8qdgf4n4g3pfee6491hc0brc4i.apps.googleusercontent.com&redirect_uri=urn%3Aietf%3Aawg%3Aoauth%3A2.0%3Aob&scope=email%20https%3A%2F%2Fwww.googleapis.com%2Fauth%2Fdocs.test%20https%3A%2F%2Fwww.googleapis.com%2Fauth%2Fdrive%20https%3A%2F%2Fwww.googleapis.com%2Fauth%2Fdrive.photos.readonly%20https%3A%2F%2Fwww.googleapis.com%2Fauth%2Fpeopleapi.readonly&response_type=code

Enter your authorization code:
.....
Mounted at /content/drive

```
In [0]: # Project Root

root_dir = ''

if USE_G_COLAB:
    root_dir = '/content/drive/My Drive/workshops/2019_07_21/sessions_01/'
```



```
In [0]: import tensorflow as tf

import requests
import numpy as np
import pandas as pd

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.utils import shuffle

import matplotlib as mpl
import matplotlib.pyplot as plt

import cv2

import random
import sys
import io
import re
import time
from datetime import datetime
import os
import struct
import itertools
from tqdm import tqdm

from pprint import pprint

%matplotlib inline
mpl.rc_file(mpl.matplotlib_fname())

%load_ext tensorboard
```

The tensorboard extension is already loaded. To reload it, use:

```
%reload_ext tensorboard
```


Persian MNIST

Data Set Information:

Attribute Information:

1. pixels

2. class:

- 0
- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9

```
In [0]: !wget https://www.dropbox.com/s/op3ht07lfou9lbz/DigitDB.zip  
        !unzip DigitDB.zip  
        !ls
```

```
--2019-07-24 14:55:36-- https://www.dropbox.com/s/op3ht07lfou9lbz/DigitDB.zip
Resolving www.dropbox.com (www.dropbox.com)... 162.125.80.1, 2620:100:6030:1::a27d:5001
Connecting to www.dropbox.com (www.dropbox.com)|162.125.80.1|:443... connected.
HTTP request sent, awaiting response... 301 Moved Permanently
Location: /s/raw/op3ht07lfou9lbz/DigitDB.zip [following]
--2019-07-24 14:55:36-- https://www.dropbox.com/s/raw/op3ht07lfou9lbz/DigitDB.zip
Reusing existing connection to www.dropbox.com:443.
HTTP request sent, awaiting response... 302 Found
Location: https://uc3d495a9fee00a812ada14c6e90.dl.dropboxusercontent.com/cd/0/inline/AlQODTFM7hxehLsMn6A_Gfpe-u2kDrnLNR9rHkQsNMDjyJJEHNV6fVIXaAupr6mGYxF1ZIE-1Q_XYmYzLEzn8iVjM3iX1eliaVOVBqHuNWtmMg/file# [following]
--2019-07-24 14:55:37-- https://uc3d495a9fee00a812ada14c6e90.dl.dropboxusercontent.com/cd/0/inline/AlQODTFM7hxehLsMn6A_Gfpe-u2kDrnLNR9rHkQsNMDjyJJEHNV6fVIXaAupr6mGYxF1ZIE-1Q_XYmYzLEzn8iVjM3iX1eliaVOVBqHuNWtmMg/file
Resolving uc3d495a9fee00a812ada14c6e90.dl.dropboxusercontent.com (uc3d495a9fee00a812ada14c6e90.dl.dropboxusercontent.com)... 162.125.80.6, 2620:100:6030:6::a27d:5006
Connecting to uc3d495a9fee00a812ada14c6e90.dl.dropboxusercontent.com (uc3d495a9fee00a812ada14c6e90.dl.dropboxusercontent.com)|162.125.80.6|:443... connected.
HTTP request sent, awaiting response... 302 FOUND
Location: /cd/0/inline2/AlQRAmjXIBvz5iF6PoNSmVOF6L4Sfjaqr2UhodiMoFvJTrYxB3WrHGUX2BQESKTgEW13DMJRvh7GV1SY1VMmygdrqCEojaGFgIffy-_1079P7RrdM7VfilljWpLMVrS-gEc1hh3FyRVhDVHfNf9Fv3yT5yiaH7NabIPUytkqb3Mr1YhSsU_NtWKp-hZKQAXCi7r6ZaT2NEWwoMYqVN1cN7b_8aSk8gy_iCVcOAVRVME4jnz0HAicTK6QHaYCFj1WKA0bUnx45J80n_-nnS4UNuAOBqCqnQlUn0nZPHfNebFo3Fj1h4xVz9jCKB6SeC3IsUbFN5cn97LDqTh21zS0tRb5/file [following]
--2019-07-24 14:55:37-- https://uc3d495a9fee00a812ada14c6e90.dl.dropboxusercontent.com/cd/0/inline2/AlQRAmjXIBvz5iF6PoNSmVOF6L4Sfjaqr2UhodiMoFvJTrYxB3WrHGUX2BQESKTgEW13DMJRvh7GV1SY1VMmygdrqCEojaGFgIffy-_1079P7RrdM7VfilljWpLMVrS-gEc1hh3FyRVhDVHfNf9Fv3yT5yiaH7NabIPUytkqb3Mr1YhSsU_NtWKp-hZKQAXCi7r6ZaT2NEWwoMYqVN1cN7b_8aSk8gy_iCVcOAVRVME4jnz0HAicTK6QHaYCFj1WKA0bUnx45J80n_-nnS4UNuAOBqCqnQlUn0nZPHfNebFo3Fj1h4xVz9jCKB6SeC3IsUbFN5cn97LDqTh21zS0tRb5/file
Reusing existing connection to uc3d495a9fee00a812ada14c6e90.dl.dropboxusercontent.com:443.
HTTP request sent, awaiting response... 200 OK
Length: 5290356 (5.0M) [application/zip]
Saving to: 'DigitDB.zip'
```

```
DigitDB.zip          100%[=====>]    5.04M   8.54MB/s    in 0.6s
```

```
2019-07-24 14:55:38 (8.54 MB/s) - 'DigitDB.zip' saved [5290356/5290356]
```

```
Archive: DigitDB.zip
  inflating: Train 60000.cdb
  inflating: RemainingSamples.cdb
  inflating: Test 20000.cdb
DigitDB.zip  RemainingSamples.cdb  'Test 20000.cdb'
drive        sample_data          'Train 60000.cdb'
```



```
In [0]: def resize_image(src_image, dst_image_height, dst_image_width):
    src_image_height = src_image.shape[0]
    src_image_width = src_image.shape[1]

    if src_image_height > dst_image_height or src_image_width > dst_image_width:
        height_scale = dst_image_height / src_image_height
        width_scale = dst_image_width / src_image_width
        scale = min(height_scale, width_scale)
        img = cv2.resize(src=src_image, dsize=(0, 0), fx=scale, fy=scale, interpolation=cv2.INTER_CUB
IC)
    else:
        img = src_image

    img_height = img.shape[0]
    img_width = img.shape[1]

    dst_image = np.zeros(shape=[dst_image_height, dst_image_width], dtype=np.uint8)

    y_offset = (dst_image_height - img_height) // 2
    x_offset = (dst_image_width - img_width) // 2

    dst_image[y_offset:y_offset + img_height, x_offset:x_offset + img_width] = img

    return dst_image

def read_cdb(filepath):
    with open(filepath, 'rb') as f:
        data = f.read()
        offset = 0

        # read private header
        yy = struct.unpack_from('H', data, offset)[0]
        offset += 2

        m = struct.unpack_from('B', data, offset)[0]
        offset += 1

        d = struct.unpack_from('B', data, offset)[0]
        offset += 1

        h = struct.unpack_from('B', data, offset)[0]
```

```
offset += 1

w = struct.unpack_from('B', data, offset)[0]
offset += 1

total_rec = struct.unpack_from('I', data, offset)[0]
offset += 4

letter_count = struct.unpack_from('128I', data, offset)
offset += 128 * 4

img_type = struct.unpack_from('B', data, offset)[0] # 0: binary, 1: gray
offset += 1

comments = struct.unpack_from('256c', data, offset)
offset += 256 * 1

reserved = struct.unpack_from('245c', data, offset)
offset += 245 * 1

if (w > 0) and (h > 0):
    normal = True
else:
    normal = False

images = []
labels = []

for i in tqdm(range(total_rec), position=0):

    start_byte = struct.unpack_from('B', data, offset)[0] # must be 0xff
    offset += 1

    label = struct.unpack_from('B', data, offset)[0]
    offset += 1

    if not normal:
        w = struct.unpack_from('B', data, offset)[0]
        offset += 1

        h = struct.unpack_from('B', data, offset)[0]
        offset += 1
```

```

byte_count = struct.unpack_from('H', data, offset)[0]
offset += 2

image = np.zeros(shape=[h, w], dtype=np.uint8)

if img_type == 0:
    # Binary
    for y in range(h):
        b_white = True
        counter = 0
        while counter < w:
            wb_count = struct.unpack_from('B', data, offset)[0]
            offset += 1

            if b_white:
                image[y, counter:counter + wb_count] = 0 # Background
            else:
                image[y, counter:counter + wb_count] = 255 # ForeGround
                b_white = not b_white # black white black white ...
            counter += wb_count
        else:
            # GrayScale mode
            data = struct.unpack_from('{}B'.format(w * h), data, offset)
            offset += w * h
            image = np.asarray(data, dtype=np.uint8).reshape([w, h]).T

    images.append(image)
    labels.append(label)

return images, labels

def load_data(datapath, img_height=32, img_width=32):
    images, labels = read_cdb(datapath)
    assert len(images) == len(labels)

    x = np.zeros(shape=[len(images), img_height, img_width], dtype=np.float32)
    y = np.zeros(shape=[len(labels)], dtype=np.int)

    for i in tqdm(range(len(images)), position=0):
        image = images[i]
        image = resize_image(src_image=image, dst_image_height=img_height, dst_image_width=img_width)
        image = image / 255

```

```

        image = np.where(image >= 0.5, 1, 0)

        x[i] = image
        y[i] = labels[i]

    x, y = shuffle(x, y, random_state=0)

    return x, y

```

Load the data

```

In [0]: trainset_dir = 'Train 60000.cdb'
        testset_dir = 'Test 20000.cdb'

        x_train, y_train = load_data(trainset_dir)
        x_test, y_test = load_data(testset_dir)

        print('Train: %s, Labels: %s' % (x_train.shape, len(y_train)))
        print('Test: %s, Labels: %s' % (x_test.shape, len(y_test)))

```

```

100%|██████████| 60000/60000 [00:07<00:00, 8401.26it/s]
100%|██████████| 60000/60000 [00:01<00:00, 58604.49it/s]
100%|██████████| 20000/20000 [00:02<00:00, 7866.61it/s]
100%|██████████| 20000/20000 [00:00<00:00, 55908.21it/s]

```

```

Train: (60000, 32, 32), Labels: #60000
Test: (20000, 32, 32), Labels: #20000

```

```

In [0]: x_train, x_valid, y_train, y_valid = train_test_split(x_train, y_train, test_size=0.05, random_state=
        RANDOM_SEED)

        print('Train: %s, Labels: %s' % (x_train.shape, len(y_train)))
        print('Valid: %s, Labels: %s' % (x_valid.shape, len(y_valid)))
        print('Test: %s, Labels: %s' % (x_test.shape, len(y_test)))

```


```

Train: (57000, 32, 32), Labels: #57000
Valid: (3000, 32, 32), Labels: #3000
Test: (20000, 32, 32), Labels: #20000


```


Visualizing the data

```
In [0]: plt.figure(figsize=(8, 8))
        for i in range(25):
            plt.subplot(5, 5, i+1)
            plt.xticks([])
            plt.yticks([])
            plt.grid(False)
            plt.imshow(x_train[i], cmap=plt.cm.binary)
            plt.xlabel(str(y_train[i]))
        plt.show()
```



```
In [0]: plt.figure(figsize=(8, 8))
        for i in range(25):
            plt.subplot(5, 5, i+1)
            plt.xticks([])
            plt.yticks([])
            plt.grid(False)
            plt.imshow(x_test[i], cmap=plt.cm.binary)
            plt.xlabel(str(y_test[i]))
        plt.show()
```



Preprocessing

In [0]:

```
Train: (57000, 32, 32, 1)
Valid: (3000, 32, 32, 1)
Test: (20000, 32, 32, 1)
```

Arch

```
In [0]: def build_model(dnn_units):
        model, r, evaluate = None, None, None

        return model, r, evaluate
```

```
In [0]: model, r, evaluate = build_model(64)
        evaluate
```

Tensorboard

Hyperparams Tuning

```
In [0]: def train_test_model(hparams):
        model = None
        accuracy = 0

        return accuracy
```

```
In [0]: from datetime import datetime
```

```
!rm -rf ./logs/  
logdir = 'logs/scaler/' + datetime.now().strftime('%Y%m%d-%H%M%S')
```

```
In [0]: from tensorboard.plugins.hparams import api as hp
```

```
HP_DNN_UNITS = hp.HParam('dnn_units', hp.Discrete([16, 32, 64]))  
HP_DROPOUT = hp.HParam('dropout', hp.RealInterval(0.1, 0.4))  
HP_OPTIMIZER = hp.HParam('optimizer', hp.Discrete(['adam', 'rmsprop', 'sgd']))  
METRIC_ACCURACY = 'accuracy'  
  
with tf.summary.create_file_writer(logdir + '/hparam').as_default():  
    hp.hparams_config(  
        hparams=[HP_DNN_UNITS, HP_DROPOUT, HP_OPTIMIZER],  
        metrics=[hp.Metric(METRIC_ACCURACY, display_name='Accuracy')],  
    )
```

```
In [0]: def run(run_dir, hparams):  
    with tf.summary.create_file_writer(run_dir).as_default():  
        hp.hparams(hparams) # record the values used in this trial  
        accuracy = train_test_model(hparams)  
        tf.summary.scalar(METRIC_ACCURACY, accuracy, step=1)
```

```
In [0]: session_num = 0
```

```
for num_units in HP_DNN_UNITS.domain.values:  
    for dropout_rate in (HP_DROPOUT.domain.min_value, HP_DROPOUT.domain.max_value):  
        for optimizer in HP_OPTIMIZER.domain.values:  
            hparams = {  
                HP_DNN_UNITS: num_units,  
                HP_DROPOUT: dropout_rate,  
                HP_OPTIMIZER: optimizer,  
            }  
            run_name = "run-%d" % session_num  
            print('--- Starting trial: %s' % run_name)  
            print({h.name: hparams[h] for h in hparams})  
            run(logdir + '/hparam/' + run_name, hparams)  
            session_num += 1
```

After Tuning

```
In [0]: def build_optimal_model():  
        model = None  
  
        return model
```

```
In [0]: optimal_model = build_optimal_model()
```

```
In [0]: from datetime import datetime  
  
!rm -rf ./logs/  
logdir = 'logs/scaler/' + datetime.now().strftime('%Y%m%d-%H%M%S')
```

```
In [0]: class ImagePredictionCallback(tf.keras.callbacks.Callback):  
  
        def __init__(self, log_dir, model):  
            super(ImagePredictionCallback, self).__init__()
```

```
In [0]: ig_cb = ImagePredictionCallback(log_dir=logdir, model=optimal_model)  
        tb_cb = tf.keras.callbacks.TensorBoard(log_dir=logdir, histogram_freq=1)
```

```
In [0]: r = optimal_model.fit(x_train_ext, y_train,  
                             validation_data=[x_valid_ext, y_valid],  
                             epochs=2,  
                             batch_size=128,  
                             callbacks=[tb_cb, ig_cb],  
                             verbose=1)
```

Tensorboard

```
In [0]: # %tensorboard --logdir logs/scaler/
```


In [0]: